

The hidden manifold distance for functional data

Susan Wei¹ and Marie-Hélène Descary²

¹ The University of Melbourne, ²Université du Québec à Montréal

Abstract

The clustering and classification of functional data, discretely-sampled curves varying over a continuum, is frequently encountered in machine learning. While one can naively treat the data as vectors in Euclidean space, one may also regard the data as functions. The challenge of analyzing infinite-dimensional objects can be ameliorated if the curves are smooth since the dimensionality is then only artificially high. This insight underpins functional data analysis, a subfield of statistics that studies functional data. For certain types of functional data such as probability density functions and warped curves of a common template function, there is further structure to exploit. Namely, certain functional data may lie on a manifold of low intrinsic dimension. For nonlinear manifolds, even basic operations such as addition and subtraction require special consideration. This work addresses the estimation of pairwise geodesic distances between noisy functional manifold data that live near, rather than on, a manifold. This setting cannot be tackled by classic manifold learning techniques which require data to live exactly on a manifold. The proposed methodology first sends the observed functional data to the hidden manifold, estimated using subspace-constrained mean-shift. Geodesic distances are subsequently calculated by employing shortest-path algorithms on this estimated manifold. Improved estimation of the pairwise geodesic distance has beneficial implications for many downstream tasks in functional data analysis which we illustrate in the case of distance-based functional classification/clustering.

Introduction

Many statistical and machine learning methods rely on some measure of distance. For example, clustering and classification of functional data, discretely-sampled data varying over a continuum, is a common machine learning task in which distance plays a crucial role. The successful development of functional data analysis, a branch of statistics dedicated to the analysis of infinite-dimensional smooth curves, demonstrates that employing the L^2 distance between the

smoothed functional observations has advantages over the Euclidean distance between functional observations treated as vectors. In this work, we advocate that for functional manifold data, i.e. functional data that lie on a manifold, the geodesic distance is potentially even better than the L^2 distance, with benefits realized in downstream tasks such as distanced-based clustering and classification of functional manifold data. In other words, when the manifold hypothesis is plausible, e.g. classes of probability density functions and classes of warped curves of a common template function, clustering and classification using the geodesic distance may give better results than using the L_2 distance.

Let X_1, \dots, X_n be a sample of n independent realizations of a random variable X that takes value in the Hilbert space $L^2([a, b], \mathbb{R})$. Suppose additionally that the function X belongs to a Riemannian manifold $\mathcal{M} \subset L^2([a, b], \mathbb{R})$ with Riemannian metric tensor g which can be used to assign a metric on the manifold as follows (Lin et al. 2014). For each point X on the manifold, the Riemannian metric tensor g has an inner product g_X on the tangent space $T_X\mathcal{M}$. The norm of a tangent vector $V \in T_X\mathcal{M}$ is defined as

$$\|V\| = \sqrt{g_X(V, V)}.$$

The geodesic distance between two functions X_i, X_j on the manifold \mathcal{M} , based on this metric tensor g , is defined as

$$d_{\mathcal{M}}(X_i, X_j) := \inf\{l(\gamma) : \gamma \text{ is piecewise-smooth function from } [a, b] \text{ to } \mathcal{M}\}$$

where

$$l(\gamma) := \int_a^b \left\| \frac{d\gamma}{dt}(t) \right\| dt$$

is the length of the curve γ . ???Is it clear how $l(\gamma)$ depends on g ???

Ideally, the functions X_i and X_j are observed on a very dense domain grid with no measurement error. Then the estimation of the geodesic distance $d_{\mathcal{M}}(X_i, X_j)$ is straightforward, e.g. a shortest-path algorithm such as Floyd-Warshall (Floyd 1962) may be used. However, most shortest-path algorithms critically assume observations live on the manifold ???add reference??? and are thus likely to fail for functional data which are almost certainly observed with noise. In this work, we put forth a technique for the recovery of the $n \times n$ matrix G of pairwise geodesic

distances, i.e.

$$G(i, j) = G(j, i) = \begin{cases} d_{\mathcal{M}}(X_i, X_j) & \text{if } i \neq j, \\ 0 & \text{otherwise.} \end{cases}$$

In lieu of X_i and X_j , we have access only to discretely-sampled noisy functional observations Y_i and Y_j that possibly lie off the true manifold.

The work of (Chen and Muller 2012) was among the first in functional data analysis to consider the manifold hypothesis for functional data. A notion of the mean and variation of functional manifold data was introduced there. Of particular relevance to this work is their proposed P-ISOMAP procedure which adds a penalty to allow for noisy functional observations whereas the classic ISOMAP algorithm (Tenenbaum, Silva, and Langford 2000) assumes observations lie exactly on the manifold. (Dimeglio et al. 2014) also realized this drawback to ISOMAP and proposed a procedure we will call robust-ISOMAP that is less sensitive to outliers. We will discuss P-ISOMAP and robust-ISOMAP in more details in the simulation section.

Proposed method for estimating geodesic distances

Suppose that each curve $X_i \in (\mathcal{M}, g)$ is observed with measurement error on a time grid $T_i = (t_{i1}, \dots, t_{iK})$, $a \leq t_{i1} < \dots < t_{iK} \leq b$, i.e. we observe a sample of K -dimensional vectors Y_1, \dots, Y_n with $Y_{ij} = X_i(t_{ij}) + \epsilon_{ij}$, where the random variables ϵ_{ij} have mean zero and are uncorrelated with each other. We assume that each observation grid T_1, \dots, T_n is dense, i.e. K is large.

We begin by converting the discretely-sampled functional observations into continuous ones. Let $\tilde{X}_1, \dots, \tilde{X}_n$ denote the functional versions of the raw data obtained by some smoothing method. For example, we may employ spline smoothing ???add reference for spline smoothing??? to recover the functional versions of the raw data, i.e.

$$\tilde{X}_i = \arg \min_{f \in C^2[0,1]} \left\{ \sum_{j=1}^K (f(t_{ij}) - Y_{ij})^2 + \lambda \|\partial_t^2 f\|_{L^2}^2 \right\}, \quad (1)$$

where $\lambda > 0$ is a tuning parameter controlling the smoothness of \tilde{X}_i . The proposed methodology for estimating the pairwise geodesic distances $\{d_{\mathcal{M}}(X_i, X_j)\}_{i>j}$ is agnostic to the smoothing method employed.

Our method is based on the idea that the underlying functional manifold \mathcal{M} can be sufficiently well-recovered by the subspace-constrained mean-shift (SCMS) algorithm (Ozertem and Erdogmus 2011). Let $\hat{M} \subset L^2([a, b], \mathbb{R})$ denote the collection of basins of attraction of the estimated probability density function of X . The SCMS algorithm sends each point $\tilde{X}_i \in L^2([a, b], \mathbb{R})$ to a point in \hat{M} . This destination is unique and we denote it $\tilde{X}_i^{\hat{M}}$. We expect the points $\tilde{X}_1^{\hat{M}}, \dots, \tilde{X}_n^{\hat{M}} \in \hat{M}$ to lie close to the real manifold \mathcal{M} . Theoretical justification that \hat{M} as estimated by SCMS is a

reasonable surrogate for the true manifold \mathcal{M} can be found in (Genovese et al. 2014).

Before presenting our procedure, we need to first review the ISOMAP algorithm (Tenenbaum, Silva, and Langford 2000), a three-step procedure that takes as input a set of points x_1, \dots, x_n in a submanifold \mathcal{M} of \mathbb{R}^D and produces an embedding in the space \mathbb{R}^d with $d < D$ that preserves pairwise geodesic distances. The procedure is as follows.

1. Construct a weighted graph G with nodes corresponding to the observations $x_1, \dots, x_n \in \mathbb{R}^D$. Two nodes x_i and x_j are connected by an edge e_{ij} with weight $d_{ij}1\{d_{ij} \leq \epsilon\}$ where $d_{ij} = \|x_i - x_j\|_2$ and $\epsilon > 0$.
2. Estimate the pairwise geodesic distances $d_{\mathcal{M}}(x_i, x_j)$ based on G using shortest-path algorithms. Specifically, the geodesic distance between two nodes is estimated to be the length of the shortest path in G between these two nodes, i.e. the sum of the weights of the edges forming the shortest path, which is calculated either with the Floyd-Warshall algorithm [Floyd1962] or with the Dijkstra algorithm [Dijkstra59anote].
3. Use multidimensional scaling ???add reference??? to obtain an embedding in \mathbb{R}^d that preserves the pairwise geodesic distances estimated above.

Note that shortest-path algorithms such as the Floyd-Warshall algorithm are used to find the smallest path *given* a weighted graph but they do not produce the weighted graphs themselves. Both P-ISOMAP and robust-ISOMAP are modifications of ISOMAP in the sense that they change the construction of the weighted graph in step 1.

In what follows, we call IsoGeo the two-step procedure which performs a modified version of the first step of ISOMAP followed by the original second step of ISOMAP. The first step of IsoGeo constructs the weighted graph G according to Dimeglio et al. (2014) so that graph is independent of the tuning parameter ϵ . Specifically, we begin by constructing the complete ???explain what complete means??? weighted graph G_c with nodes corresponding to the observations $x_1, \dots, x_n \in \mathbb{R}^D$. Next, we obtain the minimal spanning tree associated with G_c and denote its set of edges by E_s . The graph G is obtained by adding all edges e_{ij} to E_s for which the following condition is true:

$$\overline{x_i x_j} \subset \bigcup_{i=1}^n B(x_i, \epsilon_i),$$

where $B(x_i, \epsilon_i)$ is the open ball of center x_i and radius $\epsilon_i = \max_{e_{ij} \in E_s} d_{ij}$, and $\overline{x_i x_j} = \{x \in \mathbb{R}^D \mid \exists \lambda \in [0, 1], x = \lambda x_i + (1 - \lambda)x_j\}$. Let the distance estimated by IsoGeo be denoted $\text{IsoGeo}(x_i, x_j)$. Note that IsoGeo has no tuning parameters.

We are now ready to describe our procedure. Since SCMS is based on kernel density estimation, we first reduce the dimension of our data with multidimensional scaling before applying SCMS to avoid the curse of dimensionality.

1. Discretise the functions $\tilde{X}_1, \dots, \tilde{X}_n$ by evaluating them on a dense common grid $\tilde{T} = (t_1, \dots, t_D)$. Let s be a positive

integer, much smaller than D . Obtain $\tilde{X}_1^s, \dots, \tilde{X}_n^s \in \mathbb{R}^s$ using multidimensional scaling whereby the pairwise \mathbb{R}^D Euclidean distances of the discretised functions \tilde{X}_i are preserved.

2. Apply the subspace constrained mean-shift algorithm [Ozertem2011] to each of \tilde{X}_i^s to obtain $\tilde{X}_i^{s, \hat{\mathcal{M}}}$.
3. Use IsoGeo on the inputs $(\tilde{X}_1^{s, \hat{\mathcal{M}}}, \dots, \tilde{X}_n^{s, \hat{\mathcal{M}}})$ to obtain $\text{IsoGeo}(\tilde{X}_i^{s, \hat{\mathcal{M}}}, \tilde{X}_j^{s, \hat{\mathcal{M}}})_{i>j}$.

Our geodesic distance estimator is the $n \times n$ matrix \hat{G} whose elements are given by

$$\hat{G}(i, j) = \hat{G}(j, i) = \begin{cases} \text{IsoGeo}(\tilde{X}_i^{s, \hat{\mathcal{M}}}, \tilde{X}_j^{s, \hat{\mathcal{M}}}) & \text{if } i \neq j, \\ 0 & \text{otherwise.} \end{cases}$$

There are two tuning parameters to our procedure. First, there is the dimension s in step 1. In our simulations we try $s \in \{1, 2, 3\}$. The other parameter is the bandwidth h in the subspace constrained mean-shift. In our simulations we set h according to the heuristic in Equation (A1) of (Chen et al. 2015). In reality, depending on the downstream task, both s and h can be tuned in a data-adaptive way, say using cross-validation.

Simulation study

We perform a simulation study to ascertain the efficacy of our method for estimating pairwise geodesic distances. Three different metrics are used to assess the quality of a pairwise geodesic distance estimator.

The first metric assesses near- ϵ isometry, which we define as the percentage of estimated pairwise geodesic distances between $1 - \epsilon$ and $1 + \epsilon$ of the corresponding true pairwise geodesic distance. We form the receiver operating curve with ϵ on the x -axis and near- ϵ isometry on the y -axis. The metric is the area under the receiver operating curve.

The second metric $\|d - \hat{d}\|/\|\hat{d}\|$ is the Frobenius norm of the estimation error matrix $d - \hat{d}$, relative to the Frobenius norm of the true distance matrix d . The third metric is the Pearson correlation coefficient between the upper diagonal of d and \hat{d} .

Alternative estimators of geodesic distance [Marie]

We compare the performance of our method to each of the following four methods. The first method, denoted RD for raw data, takes the naive approach of applying IsoGeo directly on the raw vectors $Y_1, \dots, Y_n \in \mathbb{R}^K$. Note that this procedure is only sensible if all the grids T_1, \dots, T_n are the same. In the second method, we first smooth the discretely-sampled noisy functional data and then estimate their pairwise distance using the L^2 distance.

The third alternative, denoted SS for spline smoothing, applies IsoGeo on the smoothed version $\tilde{X}_1, \dots, \tilde{X}_n$ of the raw data obtained by spline smoothing as described in (1). The final alternative we consider is P-ISOMAP, denoted pl, which is the two-step procedure developed in Chen and Muller (2012) designed to handle noisy functional data. In the first step, a penalty is incorporated to robustify the construction of

the weighted graph. This is followed by the original second step of ISOMAP.

We ran Chen and Muller's P-ISOMAP separately in Matlab and tuned for the number of neighbors and the penalty parameter using the relative Frobenius assessment metric. We found that the penalty parameter chosen in this way was always zero, thus reducing the P-ISOMAP procedure to standard ISOMAP. Given this, we did not make further comparison to P-ISOMAP in our simulation study.

Simulation scenarios

We shall consider three functional manifolds in our simulation study. As of yet, there is little work as to how to sample from a functional manifold. Even in the Euclidean case, it is not obvious how sampling should be done (Diaconis, Holmes, and Shahshahani 2013). How to properly sample from a functional manifold could be interesting future work. For now, to safeguard against sampling unevenly on the functional manifold, we sample on a very concentrated measure for the intrinsic parameters.

Isometric functional manifold of normal density functions This scenario is modified from what is referred to as Manifold 2 in (Chen and Muller 2012) by fixing the variance of the normal density to be 1. We have

$$\mathcal{M} = \{X_\beta : \beta \in [-1, 1], t \in [a, b]\}$$

with the L_2 inner product as the metric tensor of \mathcal{M} , where $X_\beta : [a, b] \rightarrow \mathbb{R}$ is given by $X_\beta(t) = \frac{1}{\sqrt{2\pi}} \exp[-\frac{1}{2}(t - \beta)^2]$. First, note that the "straight" line connecting X_{β_1} and X_{β_2} in \mathcal{M} does not always stay inside of \mathcal{M} . Thus the geodesic distance will not coincide with the L^2 distance.

We set $a = -4$ and $b = 4$ and sample β according to a truncated standard normal with support on $[-1, 1]$. The geodesic distance between the curves X_{β_1} and X_{β_2} is given by

$$\begin{aligned} d(X_{\beta_1}, X_{\beta_2}) &= \int_{\beta_1}^{\beta_2} \left\| \frac{dX_\beta(t)}{d\beta} \right\|_{L^2} d\beta \\ &= \int_{\beta_1}^{\beta_2} \sqrt{\frac{1}{2\sqrt{\pi}} \int_{-4}^4 \frac{1}{\sqrt{\pi}} \exp\{-(t - \beta)^2\} (t - \beta)^2 dt} d\beta \\ &= \int_{\beta_1}^{\beta_2} \sqrt{\frac{1}{2\sqrt{\pi}} \int_{-4}^4 (t - \beta)^2 f(t) dt} d\beta, \text{ where } f \text{ is the cdf of } \beta \\ &\approx \int_{\beta_1}^{\beta_2} \sqrt{\frac{1}{2\sqrt{\pi}} \frac{1}{2}} d\beta \\ &= (\beta_2 - \beta_1) \frac{1}{2\pi^{1/4}}, \end{aligned}$$

where the approximation comes from the fact that we are integrating on $[a, b] = [-4, 4]$ and not on \mathbb{R} . We can see this manifold is isometric, since the geodesic distance between X_{β_1} and X_{β_2} in \mathcal{M} is proportional to $\beta_2 - \beta_1$, the Euclidean distance between the β 's.

Functional manifold of square root velocity functions

It was shown in Joshi, Srivastava, and Jermyn (2007) that the square root representation of probability density functions

gives rise to a simple closed-form geodesic distance. They consider the manifold

$$\mathcal{M} = \{\psi : [0, 1] \rightarrow \mathbb{R} : \psi \geq 0, \int_0^1 \psi^2(s) ds = 1\}$$

with the metric tensor given by the Fisher-Rao metric tensor

$$\langle v_1, v_2 \rangle = \int_0^1 v_1(s)v_2(s) ds$$

for two tangent vectors $v_1, v_2 \in T_\psi(\mathcal{M})$. Note that this coincides with the $L_2[0, 1]$ inner product. Joshi, Srivastava, and Jermyn (2007) showed that the geodesic distance between any two ψ_1 and ψ_2 in \mathcal{M} is simply

$$d_{\mathcal{M}}(\psi_1, \psi_2) = \cos^{-1} \langle \psi_1, \psi_2 \rangle.$$

We will specifically examine the square root of $Beta(\alpha, \beta)$ distributions which is supported on $[0, 1]$. That is,

$$M = \{\psi_{\alpha, \beta} : 1 \leq \alpha \leq 5, 2 \leq \beta \leq 5\}$$

where $\psi_{\alpha, \beta} : [0, 1] \rightarrow \mathbb{R}$ is the pdf of $Beta(\alpha, \beta)$. We sample α according to a truncated normal with mean 3 and variance 0.09 with support on $[1, 5]$. We sample β according to a truncated normal with mean 3.5 and variance 0.09 with support on $[2, 5]$.

Functional manifold of warping functions This is based on Equations 17 and 18 in Kneip and Ramsay (2008) but with z_{i1}, z_{i2} both set to 1. (Note that Equation 17 has a typo where the exponentials are missing negative signs.) Let $X_\alpha(t) = \mu(h_\alpha(t))$ be defined on $[-3, 3]$ where

$$\mu(t) = \exp\{(t - 1.5)^2/2\} + \exp\{(t + 1.5)^2/2\}$$

and

$$h_\alpha(t) = 6 \frac{\exp\{\alpha(t + 3)/6\} - 1}{\exp\{\alpha\} - 1}$$

if $\alpha \neq 0$ and $h_\alpha(t) = t$ otherwise.

Consider the manifold

$$M = \{X_\alpha : -1 \leq \alpha \leq 1\}.$$

We sample α according to a truncated standard normal with support on $[-1, 1]$. The geodesic distance is then

$$d(X_{\alpha_1}, X_{\alpha_2}) = \int_{\alpha_1}^{\alpha_2} \left\| \frac{dX_\alpha(t)}{d\alpha} \right\|_{L_2} d\alpha$$

which is estimated using numerical integration.

Results

The tuning parameters for the proposed method were chosen as described at the end of Section . In step 1 of our procedure, we take $D = 100$ for the size of the common time grid. We will investigate how two parameters affect the performance of the various methods in consideration. The first one is the signal-to-noise ratio ???describe??? which we fix at either 0.1 (low) or 0.5 (high). The second parameter is K , the dimension of the dscretly-observed noisy Y_i 's. We set K to either 30 (sparse) or 100 (abundant).

We show the performance of the SS and L2 method versus our method for three different values of s . We omit the result of the square root velocity scenario because all methods perform comparably. We can see from Figures 1 and 2 that ???.




Figure 1: Normal density




Figure 2: Warping function

Distance-based functional classification [Marie]

In this section, we explore whether our geodesic distance estimator has benefits for downstream analysis task. There are many tasks we could consider here such as distance-based nonparametric regression and distanced-based functional clustering, but we will focus on distance-based functional classification. It must be noted that while curve alignment, also known as curve registration, is necessarily performed as a preprocessing technique prior to clustering and classification, our geodesic distance estimator allows one to forsake this step.

For simplicity, assume the task is binary classification. Associated to each functional object x is a binary y indicating class membership. Consider the classifier proposed in Ferraty and Vieu (2003,2006) which is a functional version of the Nadaraya-Watson kernel estimator of class membership probabilities:

$$\hat{p}(y = 0|x) = \frac{\sum_{i=1}^n K[h^{-1}d(x, x_i)]1(y_i = 0)}{\sum_{i=1}^n K[h^{-1}d(x, x_i)]}$$

We shall compare our method to using L_2 distance, possibly weighted, and with curve registration already accomplished. Describe alternative methods in detail.

The bandwidth in the classifier should be tuned individually for each method. Also we might need to tune MDS dimension s since in real data, the dimension of the manifold might be much higher than encountered in the simulation scenarios where it never goes above 2.

Datasets used by functional classification papers

- Wheat, rainfall and phoneme in Aurore's paper "Achieving near-perfect classification for functional data"
- Berkeley growth curves in [ChenReiss2014].
- Tecator and phoneme in [Galeano2015] Mahalanobis technometrics paper.
- yeast cell cycle gene expression (can't find this publicly) in [LengMuller2005] "Classification using functional data analysis for temporal gene expression data"

Datasets used in functional manifold papers

- Berkeley growth, yeast cell cycle gene expression (can't find this publicly) in [ChenMuller2012]
- Tecator in [LinYao2017] contamination paper
- Berkeley growth, gait cycle in [Dimeglio2014] robust isomap paper

References

Chen, Dong, and Hans-Georg Muller. 2012. "NONLINEAR Manifold Representations for Functional Data." *The Annals of Statistics* 40 (1). Institute of Mathematical Statistics: 1–29. <http://www.jstor.org/stable/41713625>.

Chen, Y., S. Ho, P. E. Freeman, C. R. Genovese, and L. Wasserman. 2015. "Cosmic Web Reconstruction Through

Density Ridges: Method and Algorithm.” *Monthly Notices of the Royal Astronomical Society* 454 (1): 1140–56. doi:10.1093/mnras/stv1996.

Diaconis, Persi, Susan Holmes, and Mehrdad Shahshahani. 2013. “Sampling from a Manifold.” In *Advances in Modern Statistical Theory and Applications: A Festschrift in Honor of Morris L. Eaton*, Volume 10:102–25. Collections. Beachwood, Ohio, USA: Institute of Mathematical Statistics. doi:10.1214/12-IMSCOLL1006.

Dimeglio, Chloe, Santiago Gallon, Jean-Michel Loubes, and Elie Maza. 2014. “A Robust Algorithm for Template Curve Estimation Based on Manifold Embedding.” *Comput. Stat. Data Anal.* 70 (February). Amsterdam, The Netherlands, The Netherlands: Elsevier Science Publishers B. V.: 373–86. doi:10.1016/j.csda.2013.09.030.

Floyd, Robert W. 1962. “Algorithm 97: Shortest Path.” *Commun. ACM* 5 (6). New York, NY, USA: ACM: 345. doi:10.1145/367766.368168.

Genovese, Christopher R., Marco Perone-Pacifico, Isabella Verdinelli, and Larry Wasserman. 2014. “NONPARAMETRIC Ridge Estimation.” *The Annals of Statistics* 42 (4). Institute of Mathematical Statistics: 1511–45. <http://www.jstor.org/stable/43556332>.

Joshi, S., A. Srivastava, and I.H. Jermyn. 2007. “Riemannian Analysis of Probability Density Functions with Applications in Vision.” In *2007 IEEE Conference on Computer Vision and Pattern Recognition ; Proceedings*, 1664–71. Piscataway, NJ: IEEE. <http://dro.dur.ac.uk/18442/>.

Kneip, Alois, and James O Ramsay. 2008. “Combining Registration and Fitting for Functional Models.” *Journal of the American Statistical Association* 103 (483). Taylor & Francis: 1155–65. doi:10.1198/016214508000000517.

Lin, Binbin, Ji Yang, Xiaofei He, and Jieping Ye. 2014. “Geodesic Distance Function Learning via Heat Flow on Vector Fields.” In *Proceedings of the 31st International Conference on International Conference on Machine Learning - Volume 32*, II–145–II–153. ICML’14. Beijing, China: JMLR.org. <http://dl.acm.org/citation.cfm?id=3044805.3044909>.

Ozertem, Umut, and Deniz Erdogmus. 2011. “Locally Defined Principal Curves and Surfaces.” *J. Mach. Learn. Res.* 12 (July). JMLR.org: 1249–86. <http://dl.acm.org/citation.cfm?id=1953048.2021041>.

Tenenbaum, Joshua B., Vin de Silva, and John C. Langford. 2000. “A Global Geometric Framework for Nonlinear Dimensionality Reduction.” *Science* 290 (5500). American Association for the Advancement of Science: 2319–23. doi:10.1126/science.290.5500.2319.